



Lecture 17: Graphs and Unweighted Shortest Paths

Graphs show the relationship between objects/events

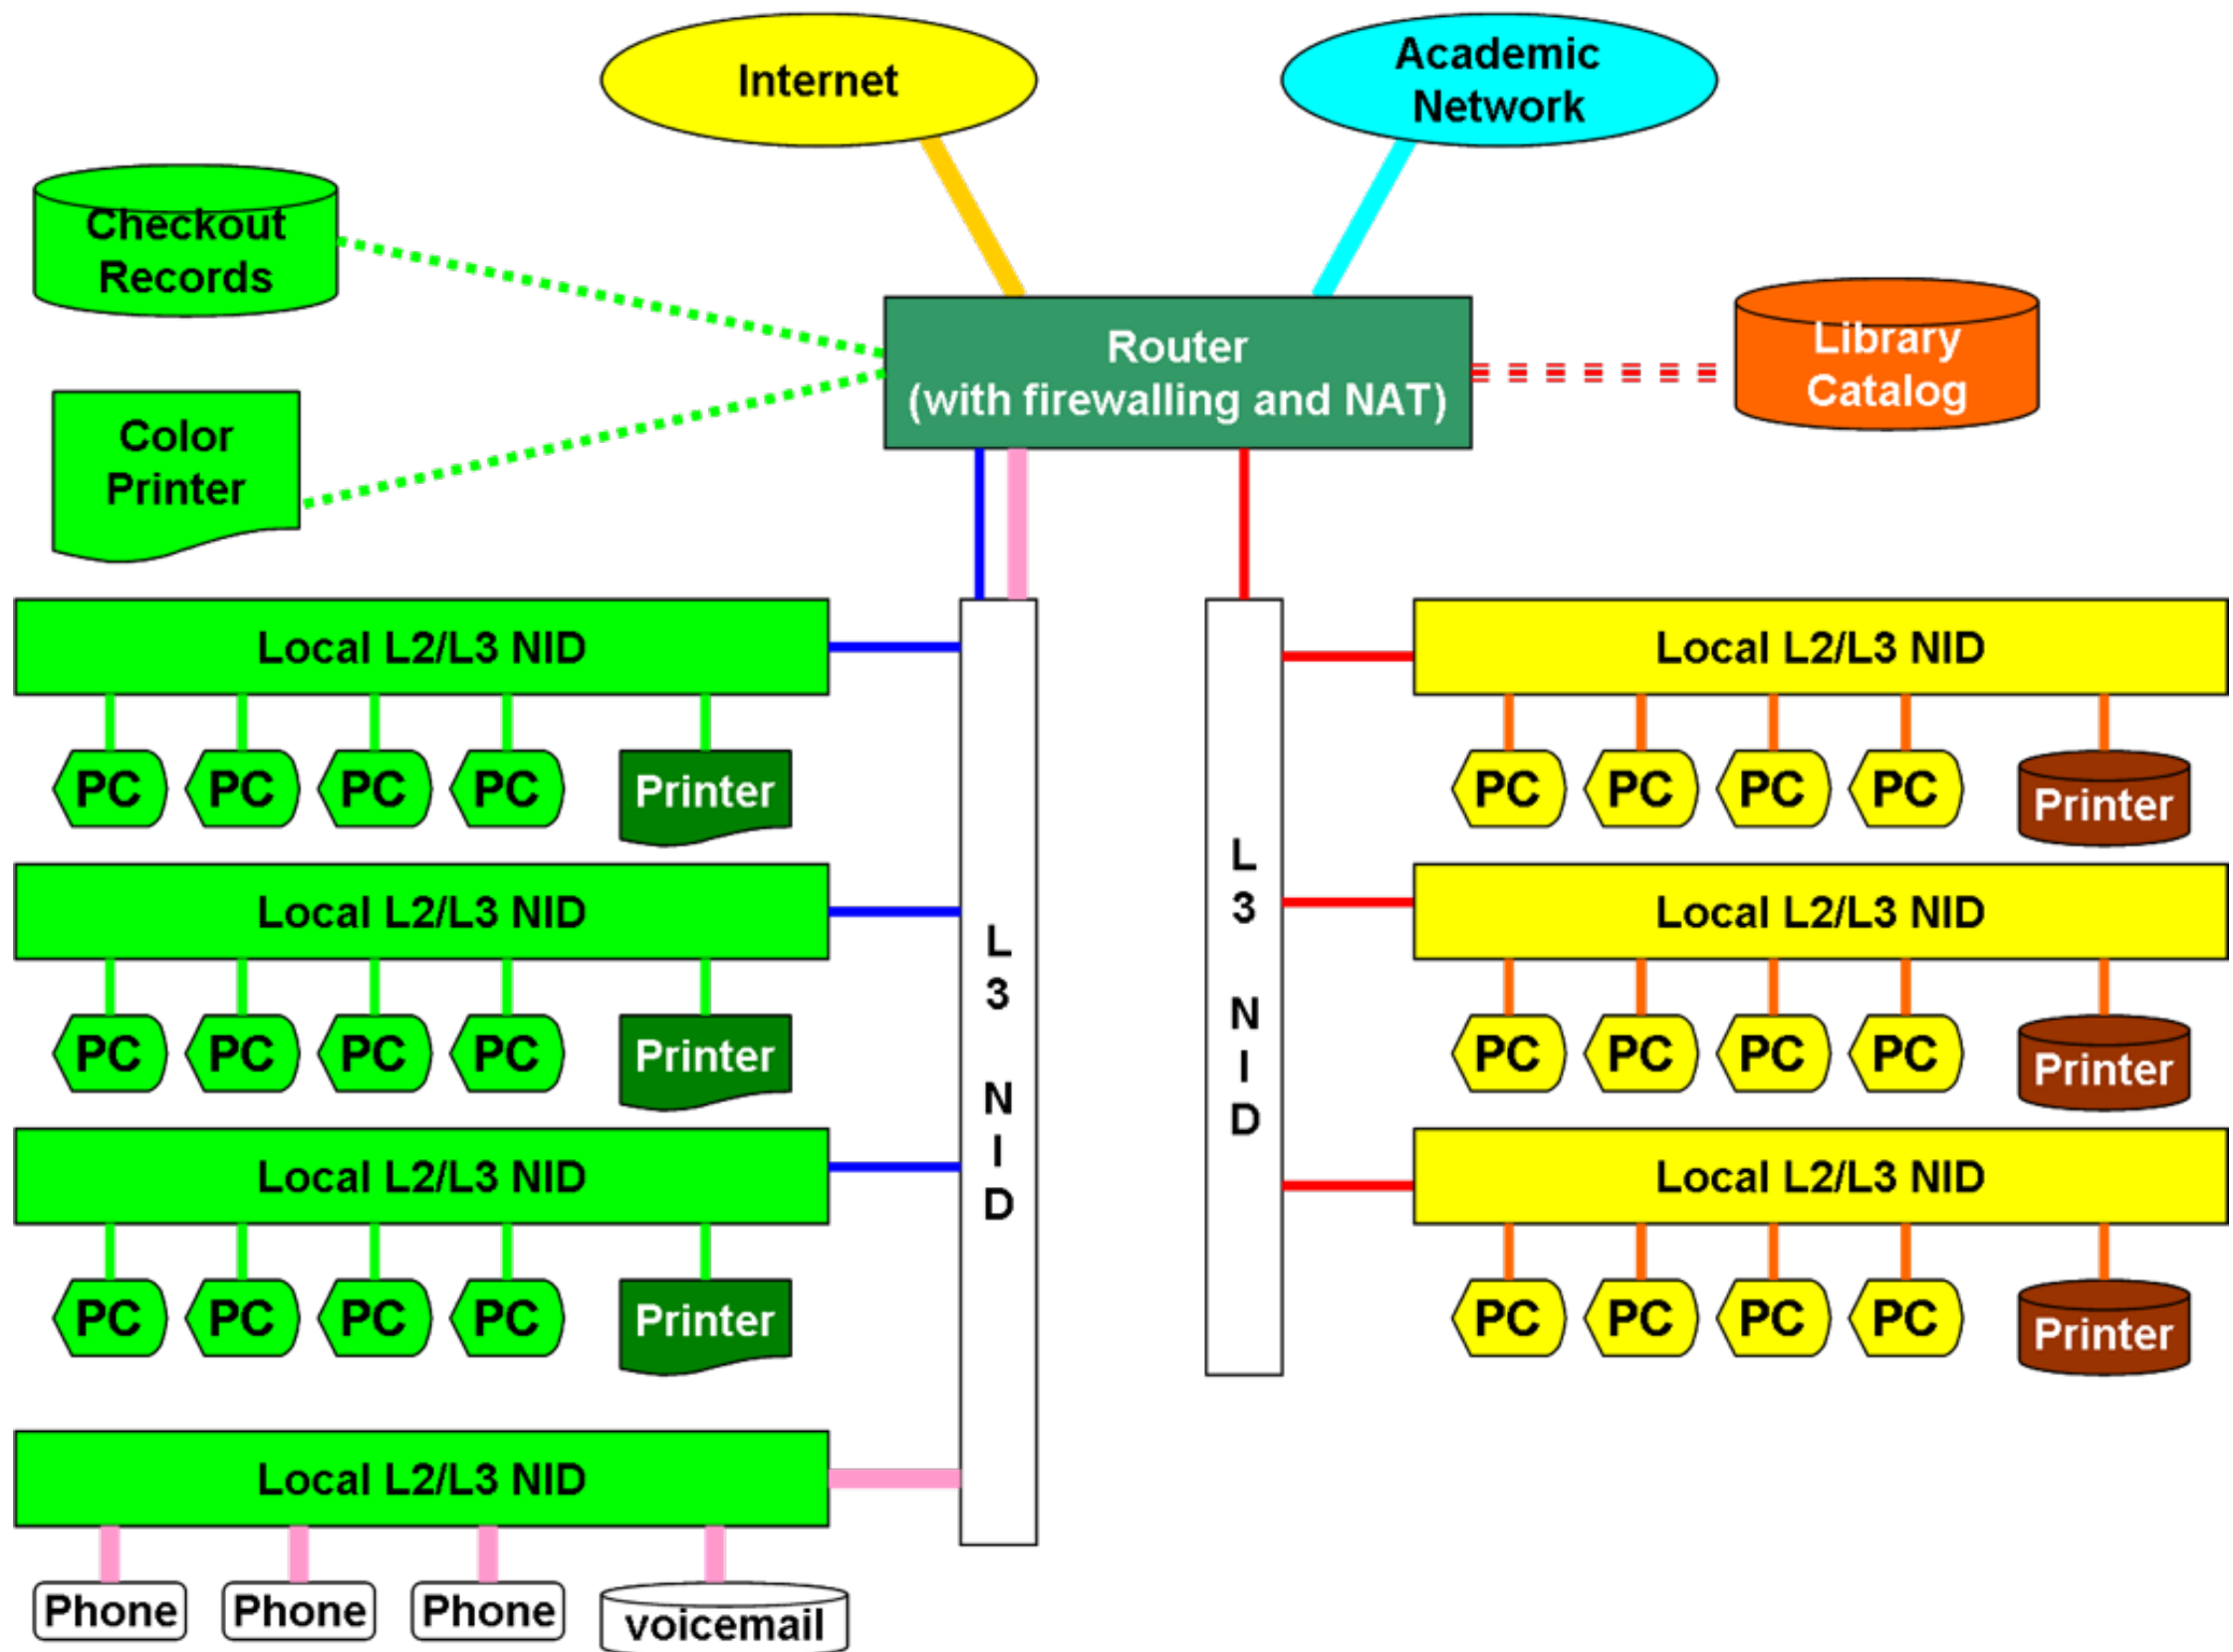
routes for mass transit



Distributed Processing



http://en.wikipedia.org/wiki/File:Distributed_Processing.jpg



<http://en.wikipedia.org/wiki/File:NETWORK-Library-LAN.png>


Graphs

- Graphs in computer science are not like the graphs you draw in physics lab.
- Mathematically, a graph $G = (V, E)$ consists of
 - a set V of vertices
 - a set E of edges
- An edge of the graph is a pair (v, w) , where v and w are vertices in V .
- Vertices are also often called nodes.
- It's often convenient to draw a picture of a graph.

Graph

- $G = (V, E)$

- V is a set of **nodes/vertices**
- E is a set of pairs of vertices, called **edges**



- Example:

- A **vertex** represent an airport

$V = \{\text{Shanghai, Tokyo, Los Angeles, Hong Kong, Sydney, Auckland, Rarotonga}\}$

- An **edge** represents a flight route between two airports

$E = \{(\text{Shanghai, Tokyo}), (\text{Shanghai, Hong Kong}), (\text{Shanghai, Sydney}),$
 $(\text{Tokyo, Hong Kong}), (\text{Tokyo, Sydney}), (\text{Tokyo, Los Angeles}),$
 $(\text{Hong Kong, Sydney}), (\text{Sydney, Auckland}), (\text{Auckland, Rarotonga}),$
 $(\text{Los Angeles, Sydney})\}$


$$|V| = 7$$

$$|E| = 10$$

A path: Los Angeles, Tokyo, Shanghai, Sydney, Auckland



This path has length 4. **Unweighted Path Length**

Graph Example




- **Undirected** graph with 4 vertices and 4 edges
- $(1,3)$ and $(3,1)$ designate the same edge

The same graph can be drawn many different ways



Directed Graph Example



Each edge has a direction, shown by an arrow





Edge (v,w) goes FROM vertex v , TO vertex w

There does **NOT** exist an edge (w,v)


Neighbor of/Adjacent to

- In an undirected graph, **w** is said to be a **neighbor** of **v** if there is an edge joining v and w. If v is a neighbor of w, then w is a neighbor of v
- In a directed graph, **w** is said to be a **neighbor** of **v** if there is a (directed) edge FROM v TO w
- “is a neighbor of” same as “is adjacent to”




Paths

- A **path** is a sequence of vertices, connected by edges.
- The **length** of the path is the number of edges.
- Example: **path** 1,2,3 has **length** 2
- Note that there may be zero, one, or more than one paths connecting two vertices.
- We often want to find the shortest path connecting two vertices.



Directed Graph:



Tokyo is adjacent to Shanghai

Shanghai is **NOT** adjacent to Tokyo
(in this graph)

$G = (V, E)$,

$V = \{\text{Shanghai, Tokyo, Los Angeles, Hong Kong, Sydney, Auckland, Rarotonga}\}$

$E = \{(\text{Shanghai, Tokyo}), (\text{Shanghai, Sydney}), (\text{Los Angeles, Tokyo}),$
 $(\text{Tokyo, Hong Kong}), (\text{Tokyo, Sydney}), (\text{Tokyo, Los Angeles}),$
 $(\text{Hong Kong, Sydney}), (\text{Sydney, Auckland}), (\text{Sydney, Hong Kong}),$
 $(\text{Sydney, Los Angeles}), (\text{Hong Kong, Shanghai}), (\text{Auckland, Sydney}),$
 $(\text{Auckland, Rarotonga}), (\text{Rarotonga, Auckland}), (\text{Los Angeles, Sydney})\}$

$|V| = 7$ $|E| = 15$


There does **NOT** exist a path: Los Angeles, Tokyo, Shanghai, Sydney, Auckland

There is a path: Los Angeles, Tokyo, Sydney, Auckland. This path has length 3.


Many applications

- **Communications:**
 - Nodes : computers
 - Edges: direct connections between computers
- **Transportation:**
 - Nodes: cities
 - Edges: roads, air routes, train tracks, etc.
- **Social**
 - Nodes: people
 - Edges: friendships, common interests, etc.

How would we store the graph ADT?



Implementation?



Adjacency Matrix:

| | Shanghai | Hong Kong | Sydney | Tokyo | Auckland | Los Angeles | Rarotonga |
|-------------|----------|-----------|--------|-------|----------|-------------|-----------|
| Shanghai | | X | X | X | | | |
| Hong Kong | X | | X | X | | | |
| Sydney | X | X | | X | X | X | |
| Tokyo | X | X | X | | | X | |
| Auckland | | | X | | | | X |
| Los Angeles | | | X | X | | | |
| Rarotonga | | | | | X | | |

Adjacency List Representation



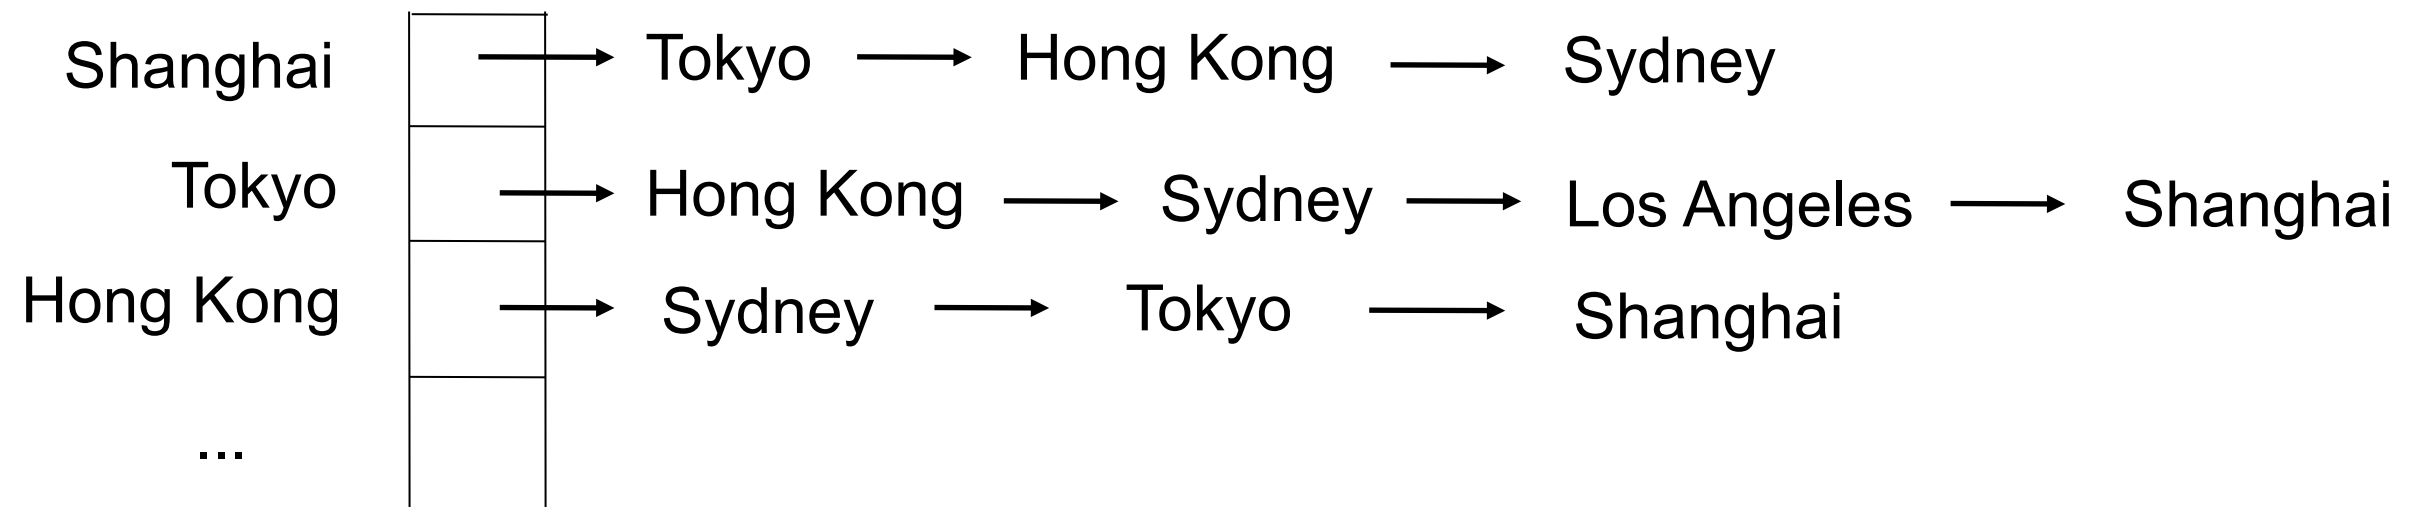


Table in position j contains list of vertices adjacent to vertex j

Adjacency List Implementation

$G = (V, E)$,


Adjacency List:



$O(|E|)$ space (assuming that every vertex is in some edge - or the number of edges is at least $O(|V|)$). **Linear in the size of the graph**

$O(|E|)$ time to construct from a list of the edges.

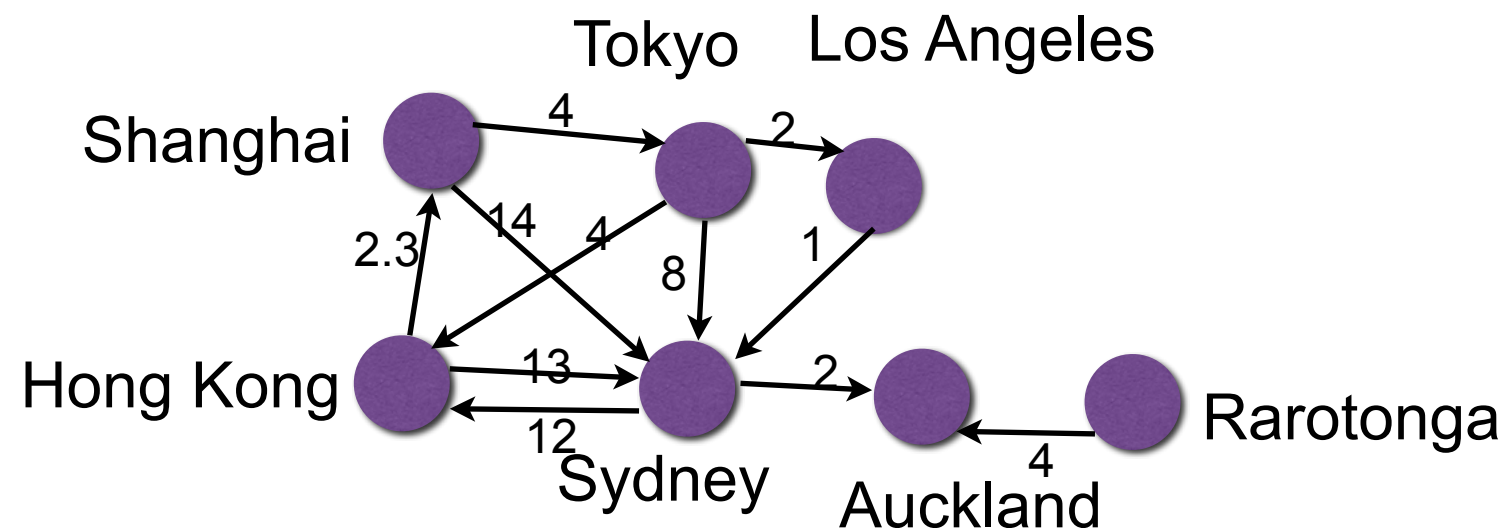
Directed Graph:



Adjacency List:

| | | | | | | |
|-----------|---|-----------|---|----------|---|-------------|
| Shanghai | → | Tokyo | → | Sydney | | |
| Tokyo | → | Hong Kong | → | Sydney | → | Los Angeles |
| Hong Kong | → | Sydney | → | Shanghai | | |
| ... | | | | | | |

Directed Graph:



Adjacency List:

| | | | | | |
|-----------|---|---------------|---|----------------|--------------------|
| Shanghai | → | Tokyo (4) | → | Sydney (14) | |
| Tokyo | → | Hong Kong (4) | → | Sydney (10) | → Los Angeles (10) |
| Hong Kong | → | Sydney (13) | → | Shanghai (2.3) | |
| ... | | | | | |

Which to Choose?

- **space**

Adjacency List: $O(V + E)$

Adjacency Matrix: $O(V^2)$

- **computational efficiency**

Adjacency List: Easy to determine all nodes adjacent to a vertex.
Easy to add an adjacent edge.

Adjacency Matrix: Easy to determine if an edge is in the graph $O(1)$ time
When not storing weights, only one bit needed per entry

- **flexibility**


Adjacency List: Cumbersome to vary the edges incident with a vertex

Adjacency Matrix: Easy to vary the edges adjacent to a vertex

Unweighted Shortest Path Problem


- Problem: Given graph **G** and source vertex **s**, find length of shortest path from **s** to every vertex in **G**.
 - **Length** of path = number of edges on the path
 - **Distance** of vertex **v** from **s** = length of shortest path from **s** to **v**

Enumerating paths and comparing their lengths would be extremely expensive: can it be done efficiently?




Finding the depth of all the nodes


Adjacency List



Use level order
search to find the
depth of each node




Unweighted Shortest-path




Breadth First Search (BFS)

- Idea:
 - Find all vertices at distance:
 - 0 from s
 - then all vertices at distance 1
 - then all vertices at distance 2
 - etc
 - Once you've found all vertices at distance k , then the vertices at distance $k+1$ are those that haven't been found yet, and that are adjacent to vertices at distance k .
 - Can implement using a queue.

Unweighted Shortest-path



Unweighted Shortest-path



Node Visiting: v2

Queue:

| | |
|----|----|
| v0 | v5 |
|----|----|

Pseudocode for Unweighted Shortest Paths

Input: Graph G , source vertex s

Let $d[v]$ denote distance from s to v

For all vertices v in G , initialize $d[v]$ to sentinel value (-1)

$d[s] = 0$

Insert s into initially empty queue Q

preprocessing

While Q is not empty {

 Delete a vertex v from front of Q (dequeue)

 For each neighbor w of v {

 If $d[w] == \text{sentinel}$ {

$d[w] = d[v] + 1$

 insert w at the back of Q (enqueue)

 }

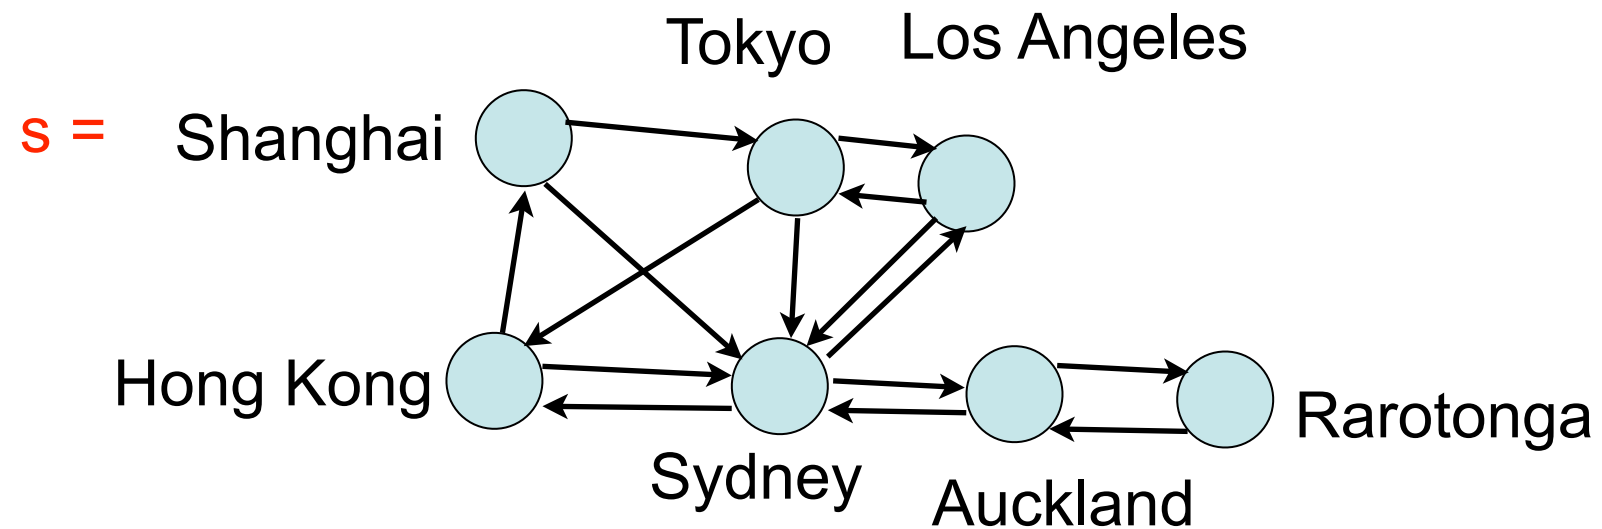
}

breadth first
search

Modification of the Pseudocode for Unweighted Shortest Paths

- Often want to know the shortest path, not just its length
 - Can modify shortest path code to also store **predecessor** on the path
 - Update when w is “discovered” as a neighbor of v .
 - At end, use these to work backwards from each target to get shortest path from source to target.
- A simple implementation is given later in the slides

Unweighted Shortest Path From Shanghai:



| | phase distance/predecessor | | | | | | | visiting | queue (front at queue on the left) |
|------|----------------------------|---|---|---|----|---|---|----------|---------------------------------------|
| | Sh | T | L | H | Sy | A | R | | |
| init | 0/- | | | | | | | | Sh |
| 1 | 0/- | | | | | | | Sh | |

Analysis of Breadth-First Search


Therefore BFS is $O(V + E)$ time - linear in the size of the adjacency-list representation of G

- It takes $O(V)$ time for queue operations
 - Every vertex/node put in queue exactly one time
 - Every vertex/node removed from the queue exactly one time
 - $O(1)$ time to push/pop/get value from a queue
- $O(V)$ time for initialization
- The adjacency list is scanned exactly one time - thus $O(E)$ time in total scanning adjacency list

// Data structures

```
typedef vector<list<int> > Graph;
// The graph is given in an adjacency list.
// Vertices are indexed from 0 to V-1
// The indices of the vertices adjacent to vertex i
// are in the list Graph[i].
// Graph can be directed or undirected.

struct vertexInf // Stores information for a vertex
{
    int dist;    // distance to vertex from the source
    int prev;    // previous node in BFS tree
};
```



```

void main()
{
    Graph g(6);

    g[0].push_back(2);
    g[1].push_back(3);
    g[1].push_back(2);
    g[2].push_back(3);
    g[2].push_back(5);
    g[2].push_back(1);
    g[2].push_back(0);
    g[3].push_back(1);
    g[3].push_back(2);
    g[3].push_back(4);
    g[5].push_back(4);
    g[4].push_back(5);
    g[4].push_back(3);
    g[5].push_back(2);
}

```

If using a singly linked list,
of course, push_front is
better - unless you are the
instructor doing this
on the whiteboard...

// Preprocessing

```
const int DEFAULT_VAL = -1;    // must be less than 0
```

```
// Breadth First Search
```

```
// The unweighted shortest path algorithm on the graph g, with vertex
```

```
// i as the source
```

```
// Prints the length (number of edges) of the shortest path from the source
```

```
// to every vertex in the graph
```

```
void shortestpaths(const Graph & g, int s)
```

```
{
```

```
    queue<int> q;                // q is the queue of vertex numbers
```

```
    vector<vertexInf> vertices(g.size()); // stores BFS info for the vertices  
                                           // info for vertex j is in position
```

```
    for (int j=0; j < vertices.size(); ++j) // Initialize distances and prev values  
        { vertices[j].dist = DEFAULT_VAL; vertices[j].prev = DEFAULT_VAL; }
```

```
    vertices[s].dist = 0;
```

```
    //rest of code on the next slide
```

```
}
```

```
//continuation of code from previous slide
```

```
q.push(s);
while (!q.empty() )
{
    int v = q.front();
    q.pop();
    for (list<int>::const_iterator w = g[v].begin(); w != g[v].end(); w++)
    {
        if (vertices[*w].dist == DEFAULT_VAL)
            //distance of *w from source not determined yet
            {
                vertices[*w].dist = vertices[v].dist+1;
                vertices[*w].prev = v;
                q.push(*w);
            }
    }
}
```

```
for (int j = 0; j < vertices.size(); j++)// print distances from source and path
{
    cout << "vertex " << j << endl;
    cout << "distance: " << vertices[j].dist << endl;
    cout << "shortest path: ";
    printpath(j,vertices);
    cout << endl;
}
```

```
}
```

```

void printpath(int j, const vector<vertexInf> & vinfo)
{
    stack<int> t;

    int current = j;
    while (current != DEFAULT_VAL)
    {
        t.push(current);
        current = vinfo[current].prev;
    }
    while (!t.empty())
    {
        cout << t.top() << " ";
        t.pop();
    }
}

```